

Hash Tables



Robert Horvick

SOFTWARE ENGINEER

@bubbafat www.roberthorvick.com





Hash Table Overview

- Associative Array

Hashing Algorithms

- Stable
- Uniform Distribution
- Secure

Hash Table Operations

- Add
- Search
- Remove

Demo: State Level Caching



Associative Array

A collection of key/value pairs where the key can only exist once in the collection



Associative Array Examples



HTTP Headers



Application
Configuration



Environment
Variables



Key/Value
Database

Request URL
Response body
Browser version
Referrer
Etc...

×	Headers	Preview	Response	Timing
▶ General				
▼ Response Headers				
age: 1062791				
alt-svc: quic=":443"; ma=604800; v="30,29,28,27,26,25"				
alternate-protocol: 443:quic,p=1				
cache-control: public, max-age=31536000				
content-length: 8056				
content-type: image/png				
date: Fri, 20 Nov 2015 15:54:55 GMT				
expires: Sat, 19 Nov 2016 15:54:55 GMT				
last-modified: Fri, 28 Aug 2015 21:45:00 GMT				
server: sffe				
status: 200				
vary: Origin				
x-content-type-options: nosniff				
x-xss-protection: 1; mode=block				
▼ Request Headers				
⚠ Provisional headers are shown				
Accept: image/webp,image/*,*/*;q=0.8				
Referer: https://www.google.com/				
User-Agent: Mozilla/5.0 (Windows NT 10.0; WOW64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/47.0.2526.73 Safari/537.36				
X-DevTools-Emulate-Network-Conditions-Client-Id: 8D870FD6-3CC3-4B4A-821E-5EA117EE89FC				



```
HTTPHeader headers;
```

```
headers["content-length"] = "8056";
```

```
headers["content-type"] = "image/png";
```

HTTP Headers

The header name and value become the key and value in the associative array



Environment Variables



```
IFS=$' \t\n'  
JAVA_HOME=/usr/local/jdk  
LESSOPEN='||/usr/bin/lesspipe.sh %s'  
LINES=50  
LOCALLIBUSER=autumnpa  
LOGNAME=autumnpa  
LS_COLORS='rs=0:di=01;34:ln=01;36:mh=00:pi=40;33:so=01;35:do=01;35:bd=40;33;01:cd=40;33;01:or=4  
:*.*arj=01;31:*.tar.gz=01;31:*.tar.bz2=01;31:*.lzh=01;31:*.lzma=01;31:*.tlz=01;31:*.txz=01;31:*.zip=01;31:*.z=01;31  
z2=01;31:*.tz=01;31:*.deb=01;31:*.rpm=01;31:*.jar=01;31:*.rar=01;31:*.ace=01;31:*.zoo=01;31:*.cp  
.pgm=01;35:*.ppm=01;35:*.tga=01;35:*.xbm=01;35:*.xpm=01;35:*.tif=01;35:*.tiff=01;35:*.png=01;35  
1;35:*.mkv=01;35:*.ogm=01;35:*.mp4=01;35:*.m4v=01;35:*.mp4v=01;35:*.vob=01;35:*.qt=01;35:*.nuv=  
v=01;35:*.gl=01;35:*.dl=01;35:*.xcf=01;35:*.xwd=01;35:*.yuv=01;35:*.cgm=01;35:*.emf=01;35:*.axv  
idi=00;36:*.mka=00;36:*.mp3=00;36:*.mpc=00;36:*.ogg=00;36:*.ra=00;36:*.wav=00;36:*.axa=00;36:*.  
LS_OPTIONS='--color=tty -F -a -b -T 0'  
MACHTYPE=x86_64-redhat-linux-gnu  
MAIL=/var/spool/mail/autumnpa  
MAILCHECK=60  
OLDPWD=/home7/autumnpa/www/drupal  
OPTERR=1  
OPTIND=1  
OSTYPE=linux-gnu  
PATH=/usr/local/jdk/bin:/usr/lib64/qt-3.3/bin:/home7/autumnpa/perl5/bin:/ramdisk/php/54/bin:/us  
home7/autumnpa/bin  
PERL5LIB=/home7/autumnpa/perl5/lib/perl5  
PERL_LOCAL_LIB_ROOT=/home7/autumnpa/perl5  
PERL_MB_OPT='--install base "/home7/autumnpa/perl5"'  
PERL_MM_OPT=INSTALL_BASE=/home7/autumnpa/perl5  
PIPESTATUS=( [0]="0")  
PPID=26704
```

```
Environment env;
```

```
env["SSH_TTY"] = "/dev/pts/0";
```

```
env["PATH"] = "/usr/local/jdk/bin/...";
```

Environment Variables

The header name and value become the key and value in the associative array



Hash Table

An associative array container that provides $O(1)$ insert, delete and search operations.



```
HashTable<String, String> headers;
```

```
headers["content-length"] = "8056";
```

```
headers["content-type"] = "image/png";
```

Hash Table

The hash table stores HTTP header data where the key and value are both strings



```
HashTable<String, HttpHeaders> headers;
```

```
headers["content-length"] = IntHeader(8056)
```

```
headers["content-type"] = StringHeader("image/png");
```

Hash Table

Hash table key and value types do not have to be the same



Hash Function

A function that maps data of arbitrary size to data of a fixed size.



Hash Function Examples



Verifying downloaded
data



Storing passwords in a
database



Hash tables key
lookup

Stability
Uniformity
Security

$$\text{hash} = f(\text{value})$$



Stability

A hash function always generates the same output given the same input.



Hash Function Stability

Stable

```
public int StableHash(string input)
{
    int result = 0;

    foreach(byte ascii in input)
    {
        result += ascii;
    }

    return result;
}
```

Unstable

```
public int UnstableHash(string input)
{
    int result = DateTime.Now.Second;

    foreach(byte ascii in input)
    {
        result += ascii;
    }

    return result;
}
```



Hash Function Stability

Stable

StableHash("foo");

StableHash("foo");

StableHash("foo");

324

324

324

Unstable

UnstableHash("foo");

UnstableHash("foo");

UnstableHash("foo");

1449447443

1449447444

1449447445



Uniformity


A hash algorithm should distribute its resulting hash value uniformly throughout the output space.




Uniform Distribution




Uniform Distribution




Uniform Distribution




Uniform Distribution




Uniform Distribution




Uniform Distribution




Uniform Distribution




Uniform Distribution



Uniform Distribution



Uniform Distribution



Uniform Distribution

(More) Uniform

```
public uint SDBMHash(string input)
{
    uint hash = 0;

    foreach (byte ascii in input)
    {
        hash = hash * 65599 + ascii;
    }

    return hash;
}
```

Non-uniform

```
public int StableHash(string input)
{
    int result = 0;

    foreach (byte ascii in input)
    {
        result += ascii;
    }

    return result;
}
```



Uniform Distribution

(More) Uniform

SDBMHash("foo");

SDBMHash("oof");

SDBMHash("of");

849955110

924308646

923718264

Non-uniform

StableHash("foo");

StableHash("oof");

StableHash("of");

324

324

324



Security


A secure hashing algorithm cannot be inverted (the input derived from the output hash).





Username	Password Hash (sdbm)
evelyn	3471203675
brian	969889485
will	1978836480





Password	Hash
aaaaaaaaa	3834880256
aaaaaaaaab	3834880257
aaaaaaaaac	3834880258
aaaaaaaaad	3834880259
...	...
zzzzzzzzz	528284160



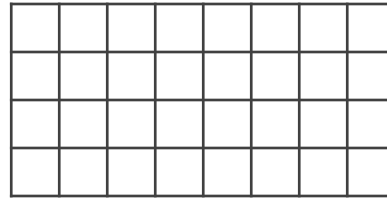


Username	Password Hash (sdbm)
evelyn	3471203675
brian	969889485
will	1978836480


Hash	Match	Actual
3471203675	aagkDhA4	password
969889485	aacOxWRs	football
1978836480	aaabhqCy	jedi




Output Size (32)




Output Size (64)




Output Size (128)



Output Size (256)



Output Size (512)



Hash Output Sizes

2^{32} (4294967296)

Checks/sec	Time
1000	49 Days
10,000	4.9 Days
100,000	12 hours
1,000,000	1 hour
10,000,000	7 minutes
100,000,000	42 seconds
1,000,000,000	4 seconds

2^{512} (1.340781e+154)

Checks/sec	Time
1000	4.25e+143 years
10,000	4.25e+142 years
100,000	4.25e+141 years
1,000,000	4.25e+140 years
10,000,000	4.25e+139 years
100,000,000	4.25e+138 years
1,000,000,000	4.25e+137 years



Sample Hash Algorithms



Additive Hash

Pros

Stable
Fast

Cons

Poor uniformity
Poor security

f	o	o
↓	↓	↓
102	111	111
102	213	324

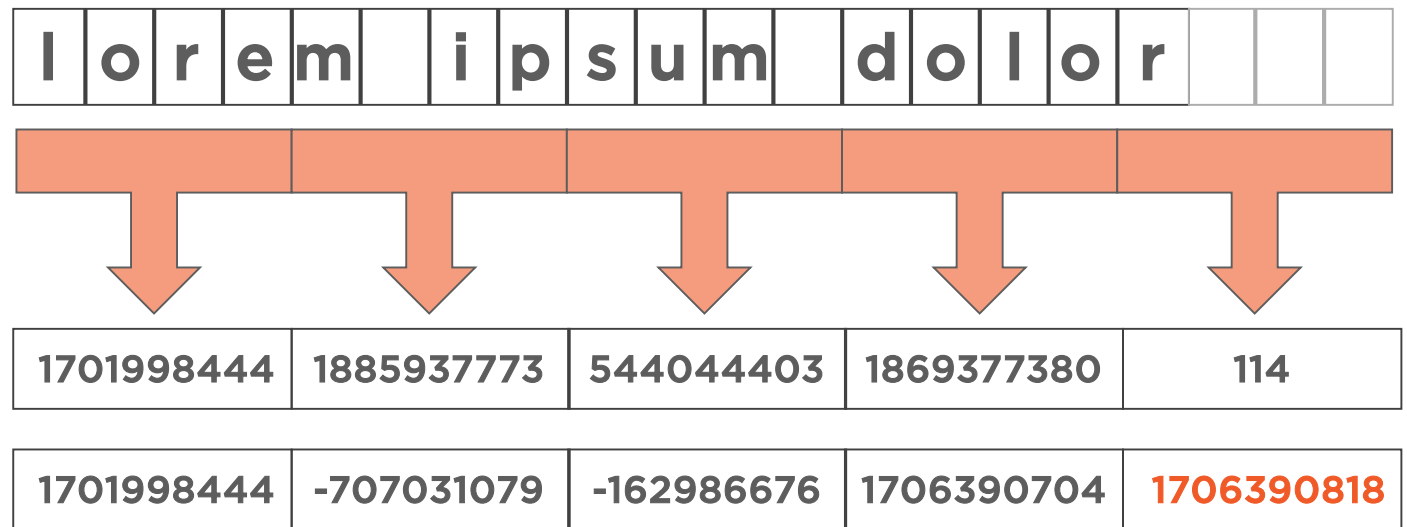
Folding Hash

Pros

- Stable
- Fast
- Better Uniformity

Cons

- Poor Security




Dbj2 Hash

```
public ulong Dbj2Hash(string input)
{
    ulong hash = 5381;

    foreach(byte c in input)
    {
        hash = hash * 33 + c;
    }

    return hash;
}
```



Hash Function Comparison

Name	Output Size	Stable	Uniform	Secure
Additive	32	YES	NO	NO
Folding	32	YES	YES	NO
Dbj2	64	YES	YES	NO
MD5	128	YES	YES	NO*
SHA-1	160	YES	YES	NO*
SHA-2	224/384	YES	YES	NO*
SHA-2	256-512	YES	YES	YES

**Once considered secure, this hash should no longer be used for secure applications.*



Adding Items



```
public class HashTable<TKey, TValue>
{
    TValue[] table = new TValue[4];

    private uint Hash(TKey key) { ... }

    public TValue this[TKey key]
    {
        get => table[Index(key)];
        set => table[Index(key)] = value;
    }

    private uint Index(TKey key)
    {
        return Hash(key) % table.Length;
    }
}
```

- ◀ Key and value type parameters
- ◀ Backing array defaults to size of 4
- ◀ A private hash function
- ◀ Functions to support indexed get and set operations using the key type
- ◀ Retrieve the value with that key
- ◀ Set the value with that key



```
HashTable<string, string> table = new HashTable<string, string>();
```

```
table["content-length"] = "8056";
```

```
table["content-type"] = "image/png";
```

Adding Values




```
HashTable<string, string> table = new HashTable<string, string>();
```

```
table["content-length"] = "8056";
```

```
table["content-type"] = "image/png";
```

Adding Values




```
HashTable<string, string> table = new HashTable<string, string>();
```

```
table["content-length"] = "8056";
```

```
table["content-type"] = "image/png";
```

Adding Values



```
HashTable<string, string> table = new HashTable<string, string>();
```

```
table["content-length"] = "8056";
```

```
table["content-type"] = "image/png";
```

Adding Values




```
HashTable<string, string> table = new HashTable<string, string>();
```

```
table["content-length"] = "8056";
```

```
table["content-type"] = "image/png";
```

Adding Values



Hash Collision

When multiple distinct keys would be inserted at the same hash table index.



Separate Chaining

Collisions in a hash table are chained together into a linked list whose root node is the hash table array entry.



```
internal class HashTableEntry<TKey, TValue>
{
    public TKey Key;

    public TValue Value;

    public HashTableEntry<TKey, TValue> Next;
}
```

Separate Chaining

The hash table handles collisions by linking all the values with the same table index into a linked list of entries.



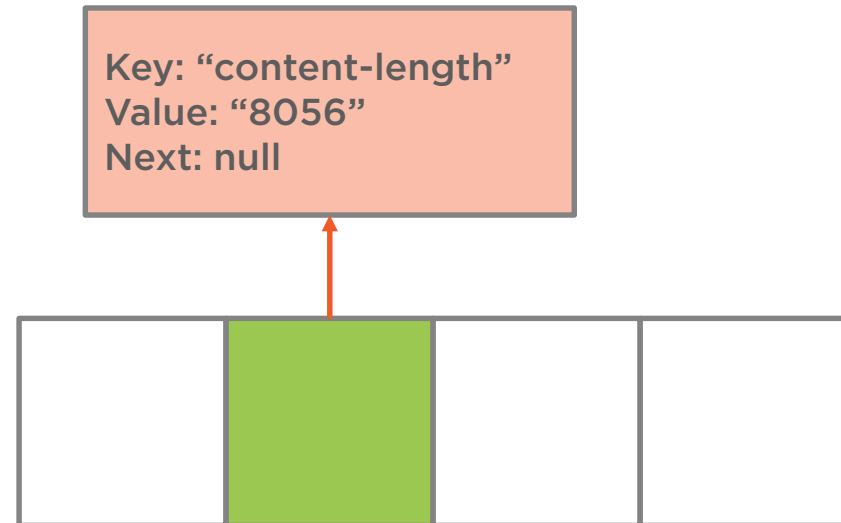
```
table["content-length"] = "8056";
```

```
table["content-type"] = "image/png";
```





```
table["content-length"] = "8056";
```

```
table["content-type"] = "image/png";
```



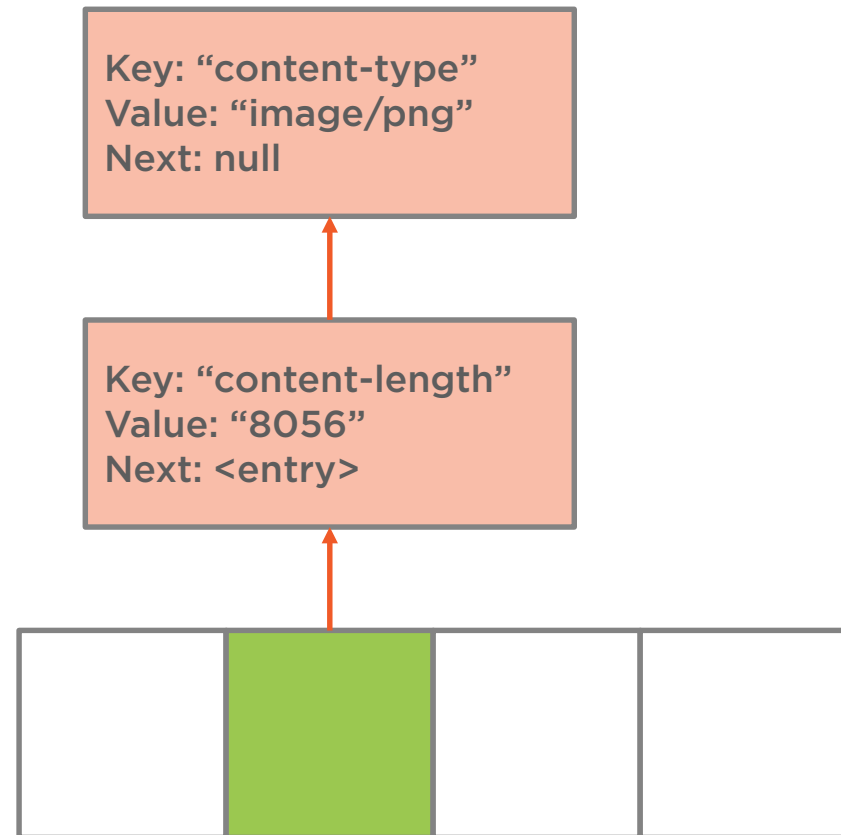
```
table["content-length"] = "8056";
```

```
table["content-type"] = "image/png";
```




```
table["content-length"] = "8056";
```

```
table["content-type"] = "image/png";
```



Hash Table Collisions



Fill Factor

The percentage of capacity representing the maximum number of entries before the table will grow. E.g., 0.80



Growth Factor


The multiple to increase the capacity of the hash table when the fill factor has been exceeded. E.g., 1.50



Hash Table Growth



Use the fill factor to determine if growth is needed



Use the growth factor to allocate a larger array




Determine the new index for the existing items in the hash table




Update the hash table to use the new array




Hash Table Growth




Hash Table Growth





Hash Table Growth





Hash Table Growth




Hash Table Growth




Hash Table Growth




Hash Table Growth




Hash Table Growth



Hash Table Growth




Hash Table Growth




Iteration






Begin at first index



Begin at first index


Visit each index



Begin at first index

Visit each index


Visit each entry




Begin at first index

Visit each index


Visit each entry




Begin at first index
Visit each index
Visit each entry
Continue to the end




Begin at first index
Visit each index
Visit each entry
Continue to the end




Begin at first index
Visit each index
Visit each entry
Continue to the end




Begin at first index
Visit each index
Visit each entry
Continue to the end



Begin at first index
Visit each index
Visit each entry
Continue to the end



Begin at first index
Visit each index
Visit each entry
Continue to the end

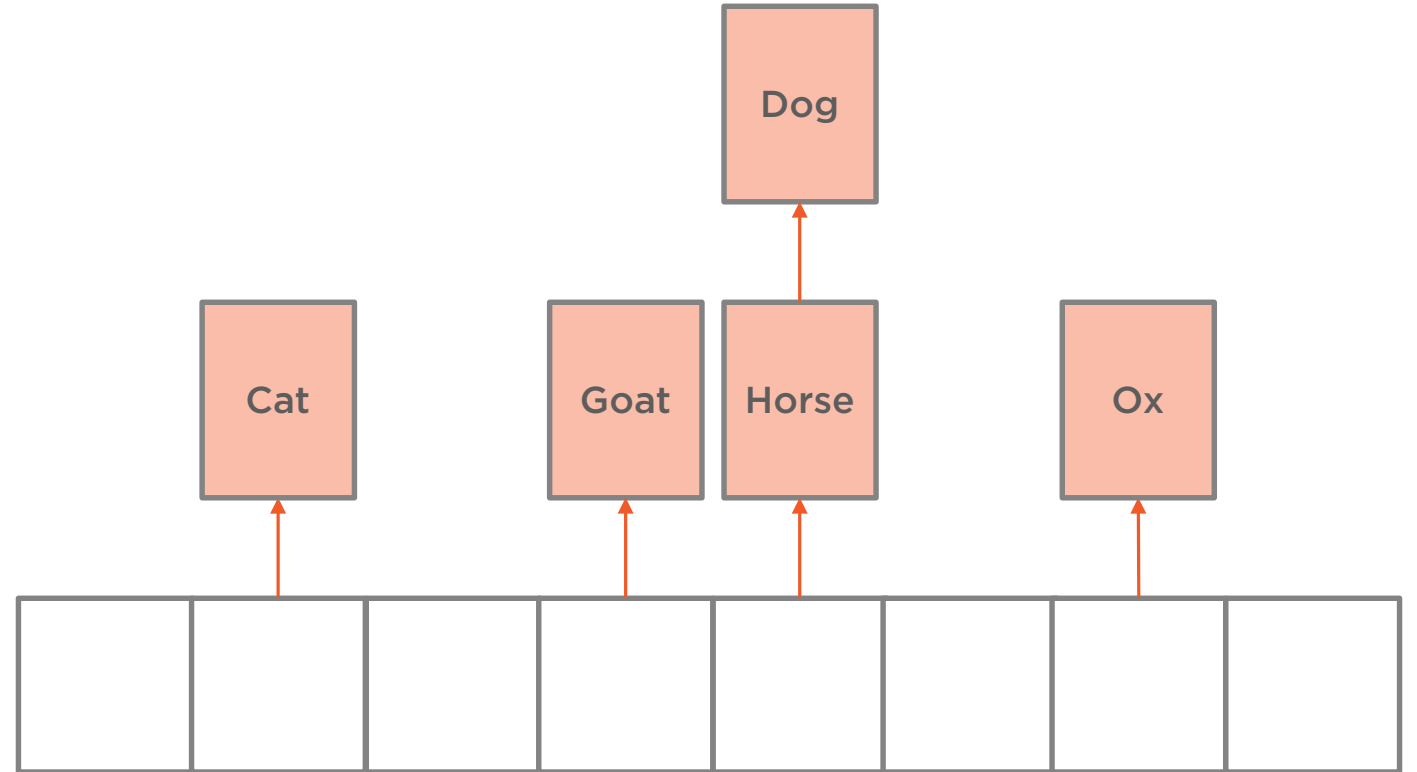


Items are iterated in the order they are stored in the table.




Finding Items



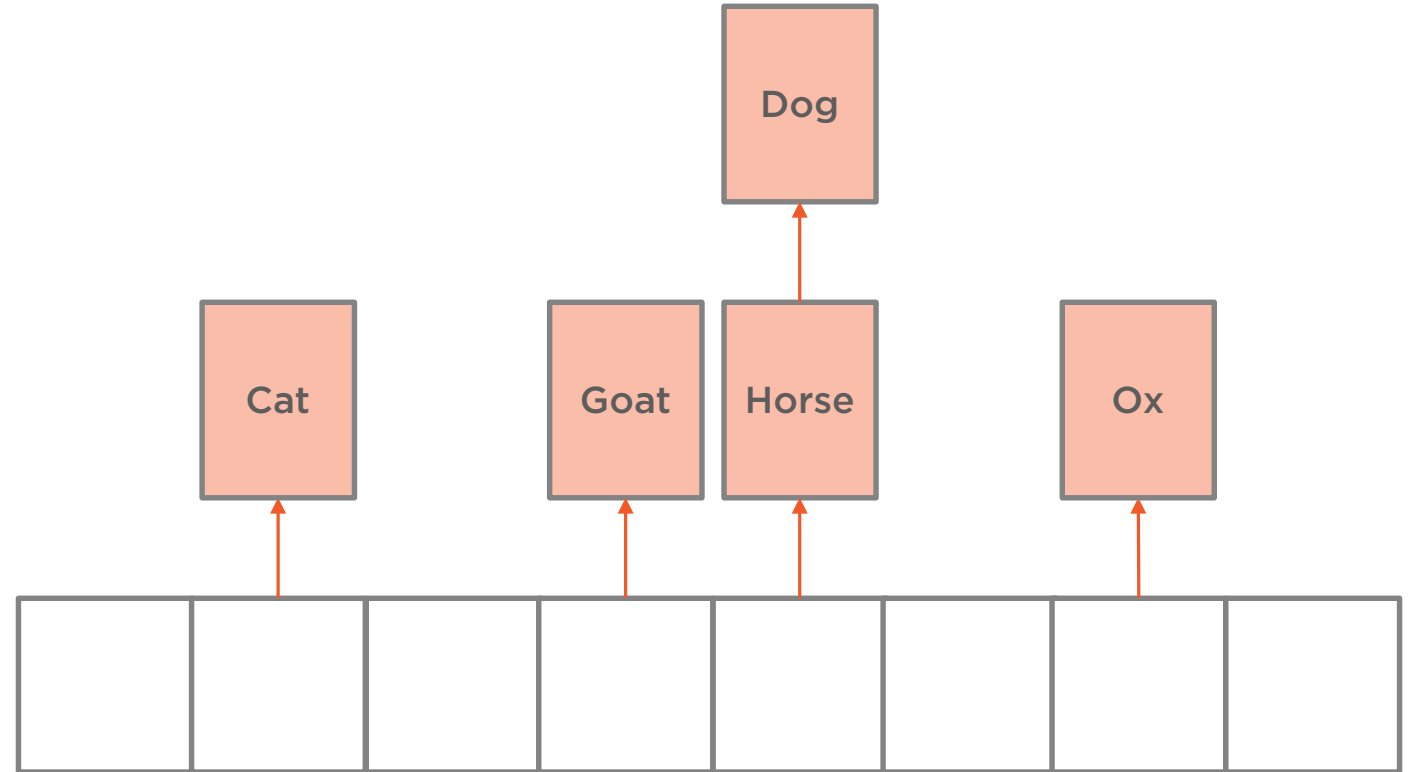


Find the index




Find the index


Use the hash function




Find the index
Use the hash function
Modulo the hash

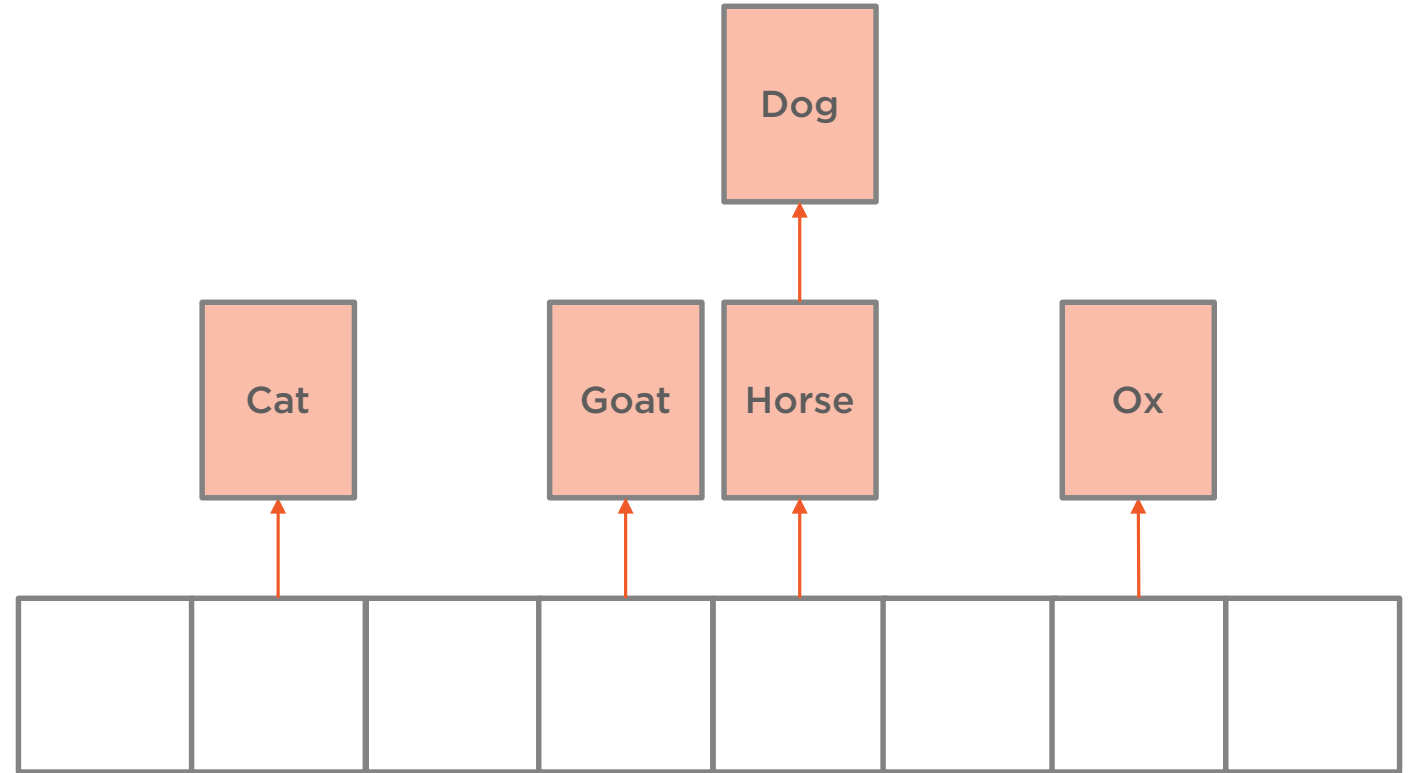


Find the index
Use the hash function
Modulo the hash
Search the entry list




Find the index
Use the hash function
Modulo the hash
Search the entry list






Find the index




Find the index

Use the hash function




Find the index
Use the hash function


Modulo the hash



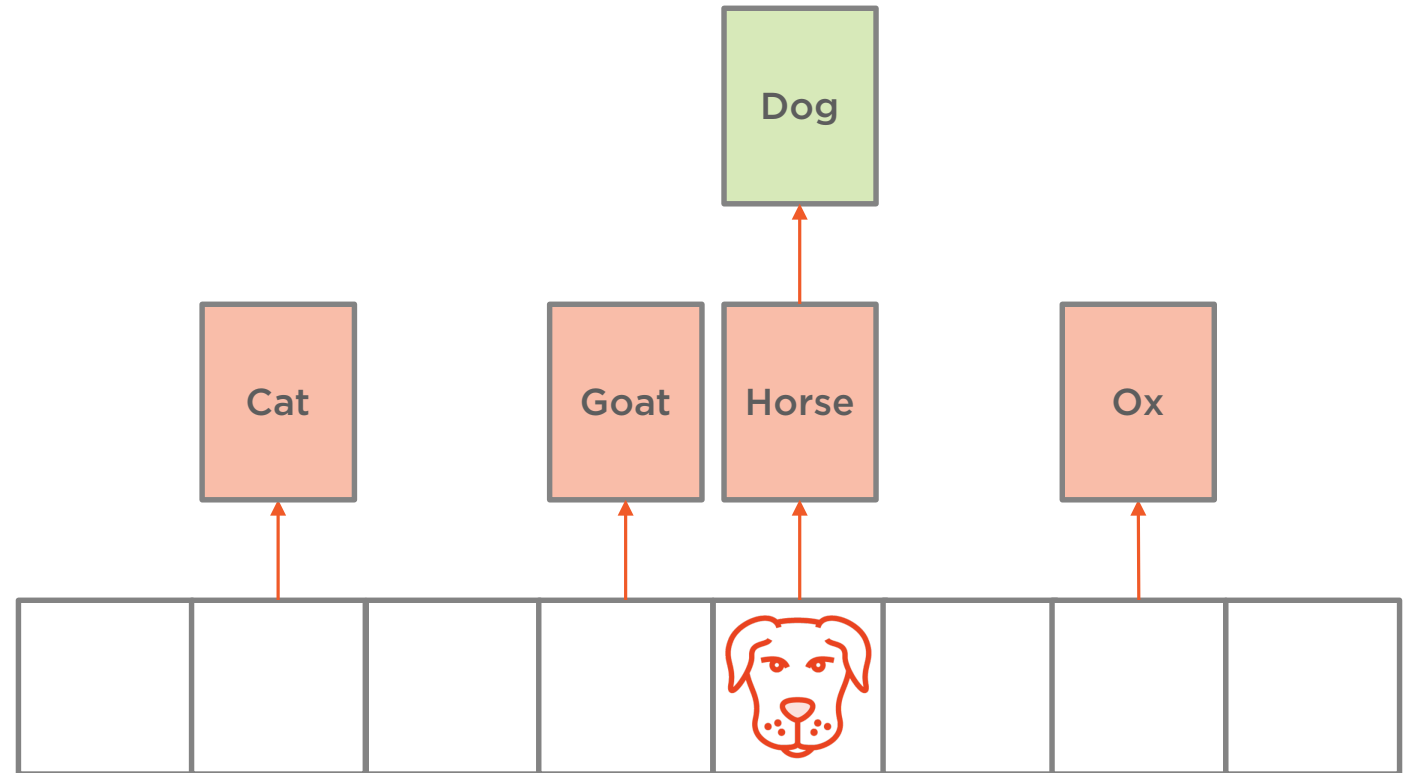
Find the index
Use the hash function
Modulo the hash
Search the entry list



Find the index
Use the hash function
Modulo the hash
Search the entry list




Find the index
Use the hash function
Modulo the hash
Search the entry list



Removing Items



Removing Items




Determine the index



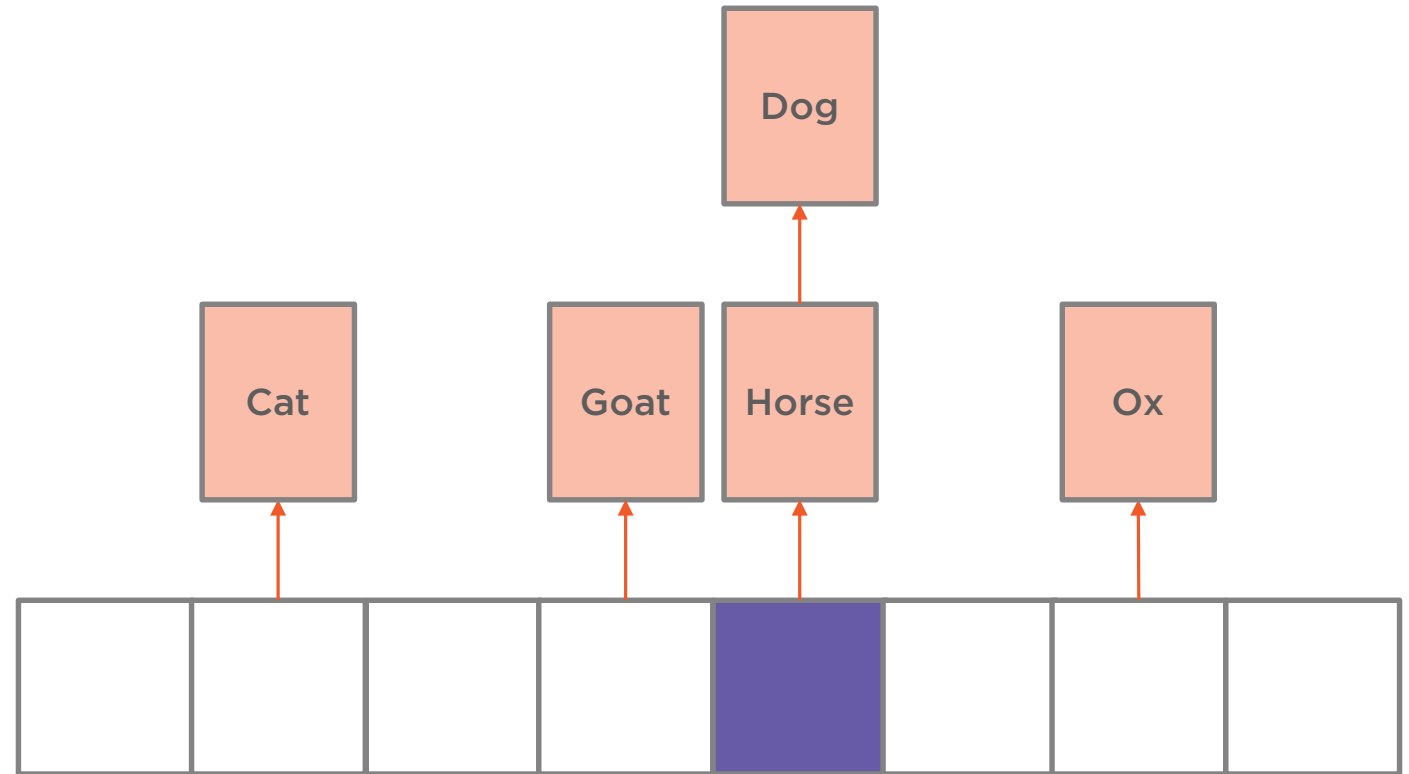
Search the entry list




Remove the entry



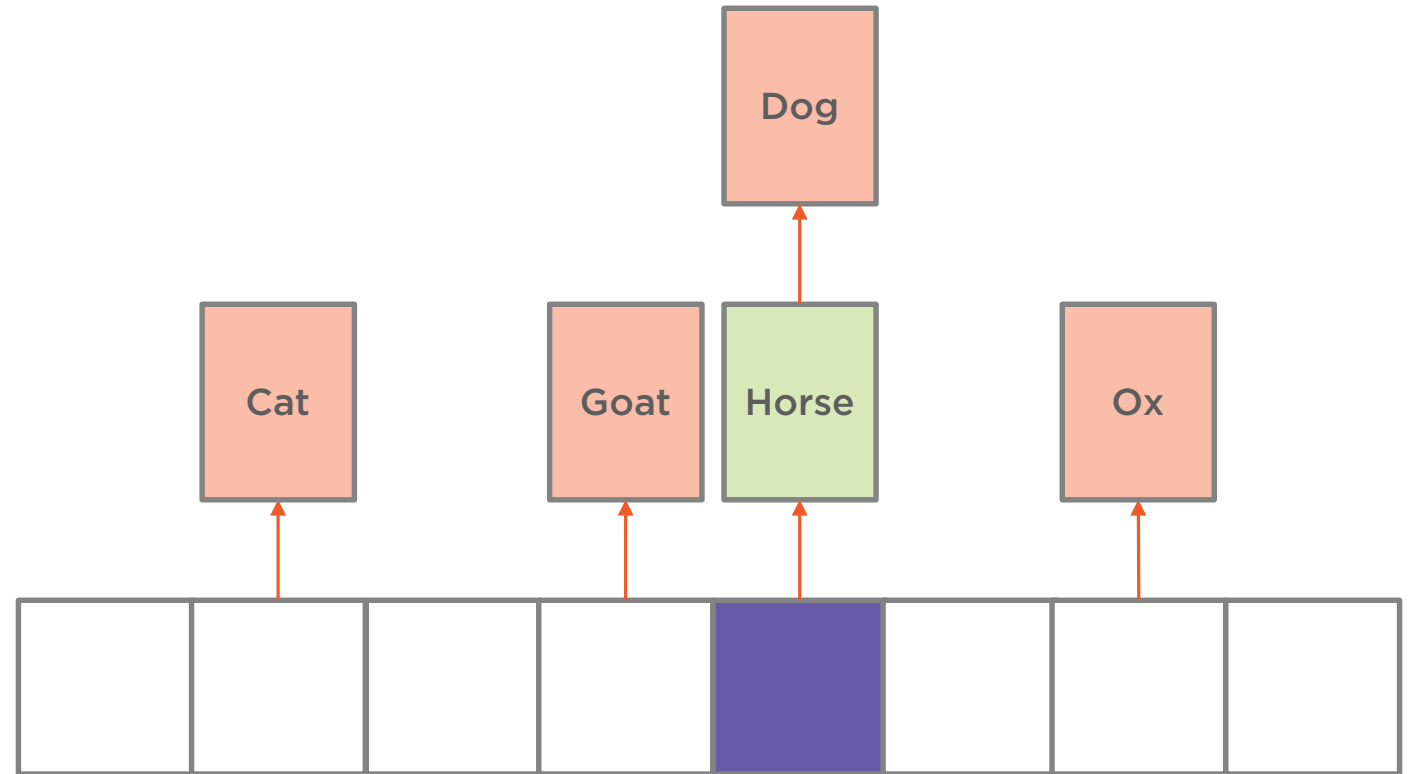
Find the index




Find the index
Search the entry list



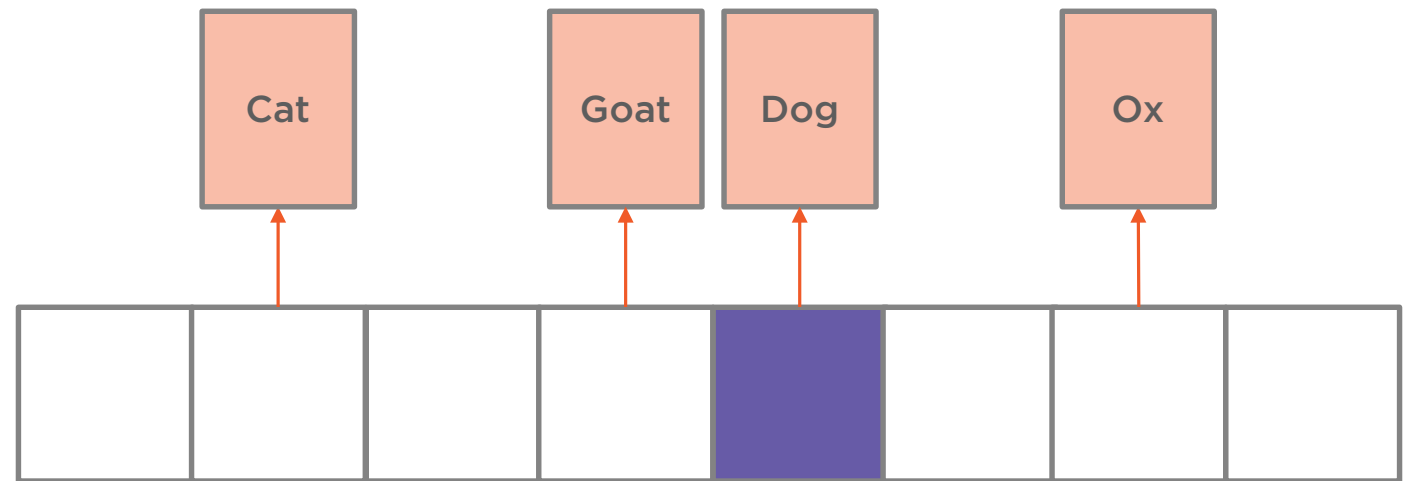
Find the index
Search the entry list
Remove the entry




Find the index
Search the entry list
Remove the entry



Find the index
Search the entry list
Remove the entry



Find the index
Search the entry list
Remove the entry



Demo



Add state-level caching

